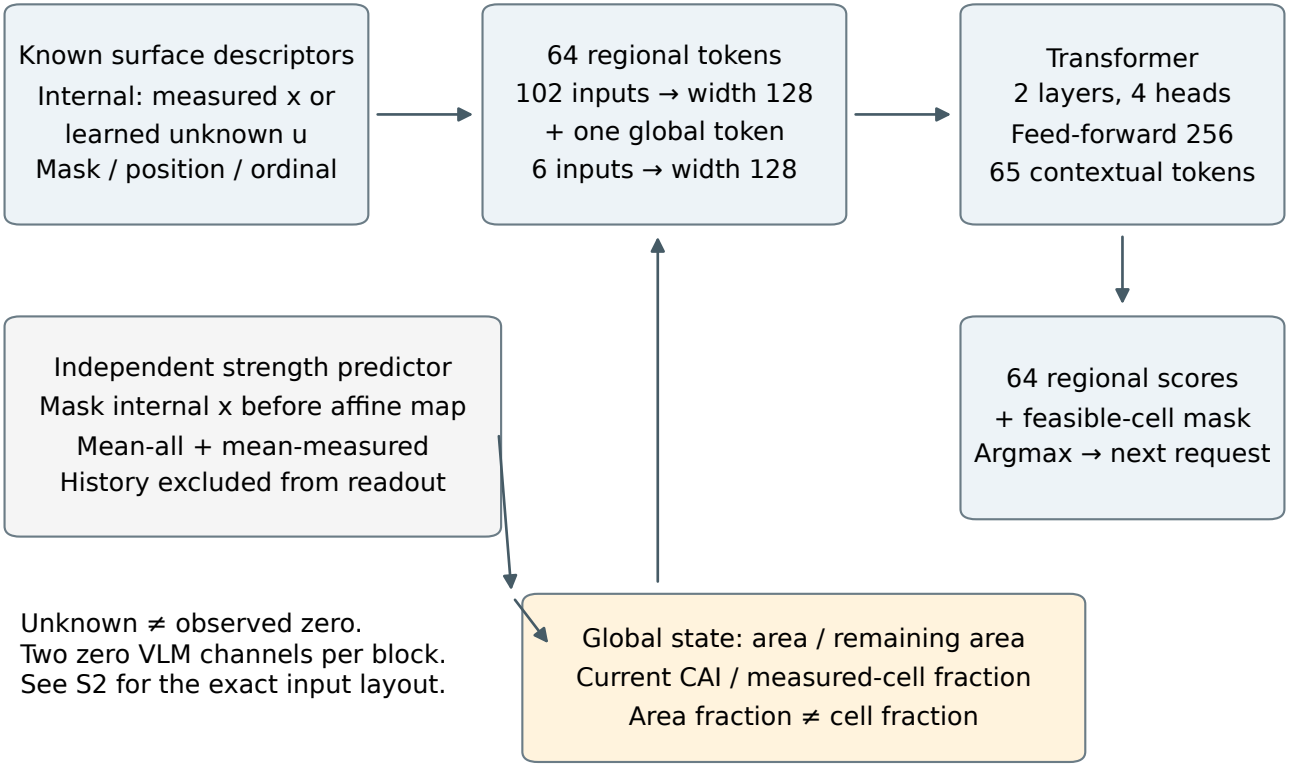
Known surface descriptors
Internal: measured x or
learned unknown u
Mask / position / ordinal
64 regional tokens
102 inputs → width 128
+ one global token
6 inputs → width 128
Transformer
2 layers, 4 heads
Feed-forward 256
65 contextual tokens
64 regional scores
+ feasible-cell mask
Argmax → next request
Independent strength predictor
Mask internal x before affine map
Mean-all + mean-measured
History excluded from readout
Global state: area / remaining area
Current CAI / measured-cell fraction
Area fraction ≠ cell fraction
Unknown ≠ observed zero.
Two zero VLM channels per block.
See S2 for the exact input layout.